

Personalized Content Discovery Using the Netflix Prize Dataset

TEAM SUBMISSION & TECHNICAL RESEARCH REPORT

Executive Abstract

This project develops and evaluates an end-to-end personalized recommendation system utilizing historical interaction records from the benchmark Netflix Prize Dataset. We design, implement, and compare three distinct recommendation approaches: an Item-Based Collaborative Filtering (IBCF) baseline, a PyTorch Latent Factor Model (SVD-like Matrix Factorization with Biases), and a deep learning Neural Collaborative Filtering (NCF) model. Operating on a dense sample of 5,000 active users and 1,500 popular movies, we analyze the trade-offs between rating prediction accuracy (RMSE) and recommendation list ranking performance (MAP@10). Our final results demonstrate that deep Neural Collaborative Filtering (NCF) achieves both the lowest rating prediction error (RMSE = 0.8213) and the highest recommendation ranking performance (MAP@10 = 0.1627) among the three models. We deploy our models via a modern, interactive Streamlit dashboard featuring explainable recommendation overlays.

Participation Format: Team (Max 02 Participants)

Institution / Platform Challenge Submission

Date: June 11, 2026

1. Problem Understanding & System Architecture

Personalized content discovery is a core driver of modern digital platforms. With billions of active users consuming media, music, and products online daily, the ability to surface highly relevant content is direct linked to platform retention, user session duration, and overall customer satisfaction.

In this study, we tackle the personalization challenge using the Netflix Prize Dataset, which was originally released in 2006 during a landmark global machine learning competition. The dataset represents one of the most sparse interaction matrices available, mirroring real-world conditions where a user interacts with less than 1% of the total content library. Our objective is twofold: (1) Rating Prediction (predicting a user's score for unseen titles) and (2) Top-K Ranking (recommending a ranked list of top movies user will like).

1.1 Recommendation Pipeline Architecture

We designed a modular recommendation pipeline containing five components:

1. Data Processing: Download, parse custom Netflix structures, sample representatively, and split data.
2. EDA Module: Compute rating and density stats, and export plots to explain catalog distributions.
3. Core Recommender Suite: House memory-based (Item-CF) and PyTorch models (SVD, NCF).
4. Evaluation Suite: Compute RMSE and MAE on test, and MAP@10 using candidate-retrieval ranking.
5. Streamlit Dashboard: Present predictions, similarities, explanations, and model comparison.

DATA FLOW FLOWCHART

Raw Netflix TXT -> Parsed DataFrame -> Sampling & User-Stratified Split (80/20)
-> Model Training Suite (Item CF / PyTorch SVD / PyTorch NCF) -> Predictions & Top-10 Recs
-> Evaluation Engine (RMSE, MAP@10 on Test Set) -> Interactive Streamlit UI / PDF Reports

2. Exploratory Data Analysis - Rating Distribution

To understand the dataset structure, we analyze the distribution of rating scores. A primary challenge in collaborative recommendation is managing user bias: users generally select items they expect to enjoy, leading to a highly positive rating bias (right-skewed distribution).




Figure 2.1: Empirical distribution of rating scores (1-5 stars) in our sampled dataset.

As visualized, 4-star ratings represent the modal class (~33.5%), followed by 3-star ratings (~28.7%) and 5-star ratings (~22.5%). Extremely poor ratings (1-star and 2-star) aggregate to less than 15.3% of the total. This positive rating bias has an important implication: a simple predictor predicting the user's average rating or global average rating (3.65) yields a respectable baseline RMSE, which latent factor models must beat by modeling specific user-item interactions.

3. Exploratory Data Analysis - Sparsity & Power Laws

A fundamental challenge in real-world personalization is dataset sparsity. In our sampled dataset containing 5,000 active users and 1,500 popular movies, the total possible interaction cells is 7,500,000. However, we observe only 1,872,548 actual interactions, corresponding to a sparsity level of 75.03% (density of 24.97%). This means collaborative filtering models must generate recommendations using only 24.97% of the matrix elements




Figure 3.1: Concentration curve illustrating movie rating volume (Lorenz curve).

Furthermore, content popularity follows a power law distribution (the long tail). As depicted in Figure 3.1, the top 20% of movies account for roughly 68.3% of the total rating interactions. This concentration indicates that platforms are dominated by blockbusters. To enhance discovery, models must be regularized properly; otherwise, they will only recommend the most popular blockbusters, neglecting the long-tail.

4. Methodology - Item-Based Collaborative Filtering

Our first recommendation model is Item-Based Collaborative Filtering (IBCF), a memory-based approach. Unlike user-based methods, which struggle because user tastes are dynamic and computation scales with user counts ($O(U^2)$), item-based collaborative filtering computes similarity between static item rating patterns. Since item count M ($1,500$) is smaller than user count U ($5,000$), the similarity matrix is compact ($M \times M$) and stable.

We implement Adjusted Cosine Similarity to compare movie rating vectors, which centers ratings by subtracting the respective user's mean rating to account for differences in rating scales (e.g. strict vs generous raters):

Adjusted Cosine Similarity Formula:

$$\text{sim}(i, j) = \frac{\sum_u (r_{u,i} - \text{mean}_u) * (r_{u,j} - \text{mean}_u)}{[\sqrt{\sum_u (r_{u,i} - \text{mean}_u)^2} * \sqrt{\sum_u (r_{u,j} - \text{mean}_u)^2}]}$$

We introduce a shrinkage factor (regularization) of 10 to penalize similarities computed on very few overlapping ratings, reducing false similarities. To predict user u 's rating for item i , we compute a similarity-weighted average of u 's ratings on items j that are in the top- K ($K=40$) most similar neighbors:

$$\text{Pred}(u, i) = \text{mean}_u + \frac{\sum_{j \text{ in neighbors}} (\text{sim}(i, j) * (r_{u,j} - \text{mean}_u))}{\sum_{j \text{ in neighbors}} |\text{sim}(i, j)|}$$

To make this computationally feasible for real-time applications, we implement a fully vectorized NumPy solver. Instead of looping over users and items, we precompute prediction scores for all cells at once via: $\text{pred} = \text{mean}_u_vec + (\text{centered_ratings} * \text{sim_top_k}) / (\text{rated_mask} * \text{sim_top_k})$ in a single matrix multiplication, executing in less than 0.5 seconds.

5. Methodology - PyTorch Matrix Factorization (SVD)

Our second model is a Latent Factor Model based on Singular Value Decomposition (SVD), implemented in PyTorch. While memory-based collaborative filtering relies on raw rating vectors, latent factor models map both users and movies to a joint low-dimensional latent space of dimension $d=50$. This compression allows the model to capture abstract characteristics (e.g. movie genres, actors, pacing, or user preferences for specific subgenres) and resolve data sparsity.

Matrix Factorization Rating Prediction:

$$\text{Pred_rating}(u, i) = \text{dot}(p_u, q_i) + b_u + b_i + \mu$$

where p_u is User Embedding, q_i is Movie Embedding, b_u/b_i are Biases, μ is Global Mean.

The model is parameterized by:


1. User Embeddings: $P \in \mathbb{R}^{U \times 50}$ and Movie Embeddings: $Q \in \mathbb{R}^{M \times 50}$.
2. User Biases: $b_u \in \mathbb{R}^U$ representing individual rating tendencies (e.g. strict vs loose raters).
3. Movie Biases: $b_i \in \mathbb{R}^M$ representing movie-specific popularity offsets.
4. Global Mean: μ , the static mean rating of the training dataset.

We optimize the parameters using gradient descent on Mean Squared Error (MSE) loss. To prevent overfitting, we add L2 regularization (weight decay = 0.01) on the embedding weights and biases. We train the network for 12 epochs using the Adam optimizer with a learning rate of 0.005 and a batch size of 512, processing ratings in mini-batches. Once trained, generating Top-K recommendations is highly scalable: we run a single forward pass multiplying the user's embedding vector p_u by the entire movie embedding matrix Q .

6. Methodology - Deep Neural Collaborative Filtering

To capture non-linear interactions, we implement Neural Collaborative Filtering (NCF), a deep learning architecture proposed by He et al. (2017). Standard matrix factorization relies on the dot product, which linearly combines latent dimensions. This linear restriction can limit recommendation performance when user-item interactions follow complex non-linear patterns.

Our NCF model combines two components in parallel: a Generalized Matrix Factorization (GMF) model and a Multi-Layer Perceptron (MLP):



1. GMF Component: Learns linear interactions by taking the element-wise product of user and item embeddings (32d).
2. MLP Component: Learns non-linear interactions. It concatenates MLP embeddings (32d), passing the 64-dimensional concatenated vector through four dense hidden layers: 64 -> 32 -> 16 -> 8. Each layer uses ReLU activation, dropout ($p=0.2$) to prevent overfitting, and Xavier initialization.

We concatenate the GMF output (32d) and MLP output (8d) into a 40-dimensional fusion vector, which is mapped via a linear regression layer to produce the rating prediction. We train NCF using PyTorch for 10 epochs, with Adam optimizer at a learning rate of 0.001.

7. Evaluation Strategy & Metrics

To measure recommendation effectiveness, we implement a rigorous validation framework:

1. Train-Test Split: We execute a User-Stratified 80/20 Split. For every user, exactly 80% of their rating history is randomly allocated to the training set, and the remaining 20% is assigned to the test set. This ensures that all test users have interaction histories in training, allowing collaborative models to form embeddings and predictions for them, avoiding cold-start user evaluation errors.

2. Rating Prediction Metrics: We measure error magnitudes on unseen test interactions using:

- RMSE (Root Mean Squared Error): Penalizes larger errors quadratically. Crucial for rating accuracy.
- MAE (Mean Absolute Error): Measures average absolute deviations.

3. Recommendation Ranking Metrics: Standard recommendation systems do not merely predict rating magnitudes; they retrieve a ranked list of relevant items. We define a movie as 'relevant' if its actual rating in the test set is greater than or equal to 3.5. We generate Top-10 recommendations from unseen items (items not in training) and evaluate:

- MAP@10 (Mean Average Precision @ 10): Measures recommendation list quality. Average Precision (AP@10) rewards models for placing relevant items higher in the list, and MAP@10 averages these scores across all test users:

$$AP@10 = \sum_{j=1..10} (\text{Precision}@j * I(\text{item}_j \text{ is relevant})) / \min(|\text{Relevant}|, 10)$$

- NDCG@10 (Normalized Discounted Cumulative Gain): Measures list utility incorporating logarithmic position discounts.

- Precision@10 & Recall@10: Measures hit proportions and coverage.

8. Experimental Results & Benchmark Analysis

We trained and evaluated all three models on our sampled Netflix dataset. The table below compiles the rating prediction errors, ranking recall/precision, and training complexity benchmark results:

Model	RMSE (L)	MAE (L)	MAP@10 (H)	NDCG@10 (H)	Precision@10	Train Time
Item-Based CF	0.8675	0.6684	0.0179	0.0355	0.0373	1.16s
Matrix Factorization	0.9131	0.7116	0.0276	0.0746	0.0799	230.15s
Neural CF (NCF)	0.8213	0.6347	0.1627	0.2840	0.2586	95.54s

Analysis of Results:

1. Rating Prediction (RMSE): Deep Neural Collaborative Filtering (NCF) achieved the lowest error (RMSE = 0.8213), outperforming Item-Based CF (RMSE = 0.8675) and Matrix Factorization (RMSE = 0.9131). This indicates that NCF's hybrid architecture (combining linear GMF and non-linear MLP layers) is highly effective at minimizing prediction residuals on the test set.
2. Ranking Performance (MAP@10): Neural CF (NCF) achieved the highest ranking quality (MAP@10 = 0.1627, NDCG@10 = 0.2840, and Precision@10 = 0.2586), far outperforming Matrix Factorization (MAP@10 = 0.0276) and Item-Based CF (MAP@10 = 0.0179). This dramatic difference highlights NCF's capacity to learn high-quality ranked lists. Its non-linear layers learn expressive decision boundaries that place relevant items at the top of the user's recommendation shelf.
3. Computational Complexity: Item-CF requires very little training time (1.16 seconds) but suffers from high computational overhead at inference time due to similarity searches. Matrix Factorization takes longer to train (230.15 seconds for 12 epochs) due to evaluation steps, while NCF trains in 95.54 seconds (10 epochs) and provides highly scalable, real-time recommendation retrieval via forward passes.

9. Deliverables, Explanations, & Future Scope

1. Explainable Recommendations:

We build explainability directly into our personalized discovery module. For any recommended movie i for user u , we scan the user's high-rated training movies ($r_{uj} \geq 3.5$), identify the movie j with the highest similarity to i in the similarity matrix, and generate an explanation:

'Recommended because you rated "Finding Nemo" a 5.0, and these movies share rating similarity (0.81).'

2. Cold Start Strategy:

Collaborative systems struggle with new items (no ratings) or new users (no history). We implement a fallback system:

- New Users: Fallback to the platform's most popular movies (highest interaction count) or highest average-rated movies.

- New Movies: Fallback to content-metadata similarity by recommending them to users who frequently rate movies from the same release year.

3. Interactive Dashboard:

We deploy our system via Streamlit, providing a premium dark-themed UI for users to explore EDA statistics, select User IDs to see their watch history and Top-10 recommendations (with NLP explanations), query movie similarity tables, and compare model performance charts.

4. Future Extensions:

To improve recommendations further, future work should incorporate rich content metadata (e.g., director, actors, genre descriptors) into a Hybrid Model using LightFM or a Factorization Machine, allowing content features to resolve the cold start problem completely.